

# ***Operating systems***

By  
I Ravindra kumar, B.Tech, M.Tech,(Ph.D.)

Assistant professor,  
Dept of CSE, VNR VJIET

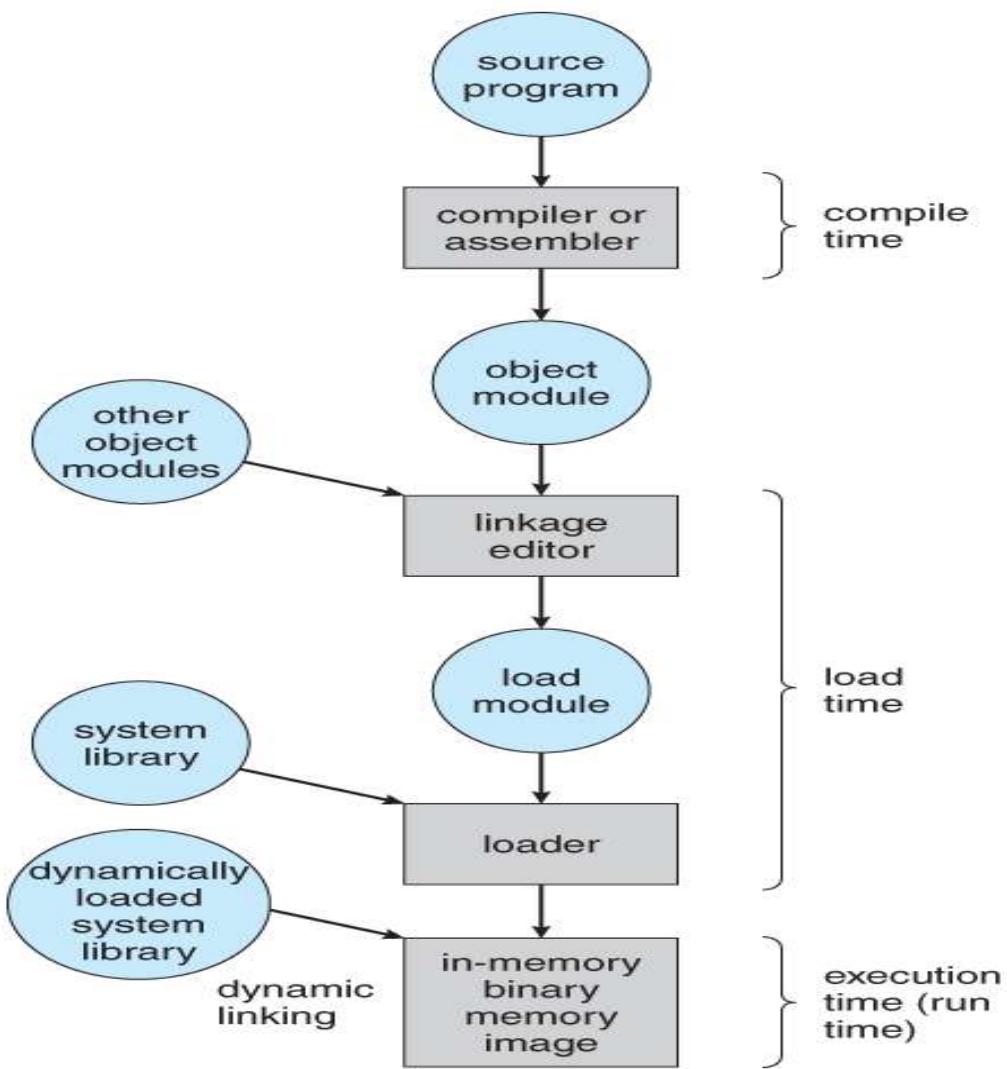
## Memory Management

## Memory Management :

- Memory consists of a large array of words or bytes, each with its own address.
- The CPU fetches instructions from memory according to the value of the program counter.
- A typical instruction-execution cycle,
- The memory unit sees only a stream of memory addresses; it does not know how they are generated
- we can ignore how a memory address is generated by a program.
  - We are interested in only the sequence of memory addresses generated by the running program.

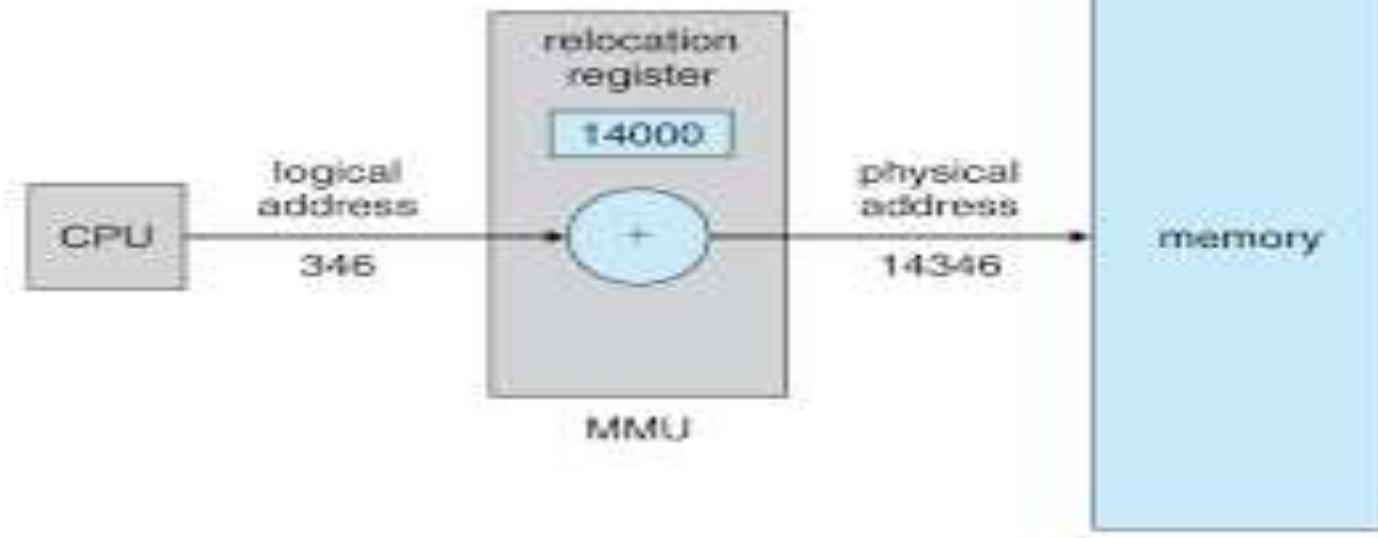
## Address Binding:

- a program resides on a disk as a binary executable file.
- The program must be brought into memory and placed within a process for it to be executed.
- The collection of processes on the disk that is waiting to be brought into memory for execution forms the input queue.
- select one of the processes in the input queue and to load that process into memory.



## Logical- Versus Physical-Address Space:

- Address generated by the CPU is commonly referred to as a logical address,
- an address seen by the memory unit—that is, the one loaded into the memory-address register of the memory—is commonly referred to as a physical address.
- usually refer to the logical address as a virtual address.
- set of all logical addresses generated by a program is a logical-address space;
- the set of all physical addresses corresponding to these logical addresses is a physical-address space.
- run-time mapping from virtual to physical addresses is done by a hardware device called the memory-management unit (MMU).
- The final location of a referenced memory address is not determined until the reference is made.



## Dynamic Loading:

- The size of a process is **limited to the size of physical memory**.
- dynamic loading, a routine is **not loaded until it is called**.
- All routines are kept on disk in a relocatable load format.
- The main program is loaded into memory and is executed.

## Dynamic Linking and Shared Libraries:

- **Static linking**, which system language **libraries** are treated like any other **object module** and are combined by the loader into the binary program image.
- **Dynamic linking**, a **stub** is included in the image for each **library-routine reference**.
- stub is a small piece of code that indicates
  - stub is executed, it **checks** to see whether the **needed routine** is already **in memory**.
  - **If not**, the program **loads** the routine into memory.
  - Either way, the **stub replaces itself** with the address of the routine, and executes the routine.
- More than one version of a library may be loaded into memory,
  - each program uses **its version information** to decide which copy of the library to use.
  - Thus, only programs that are compiled with the new library version are affected by the incompatible changes incorporated in it.
- Other programs linked before the new library was installed will continue using the older library. This system is also known as **shared libraries**.

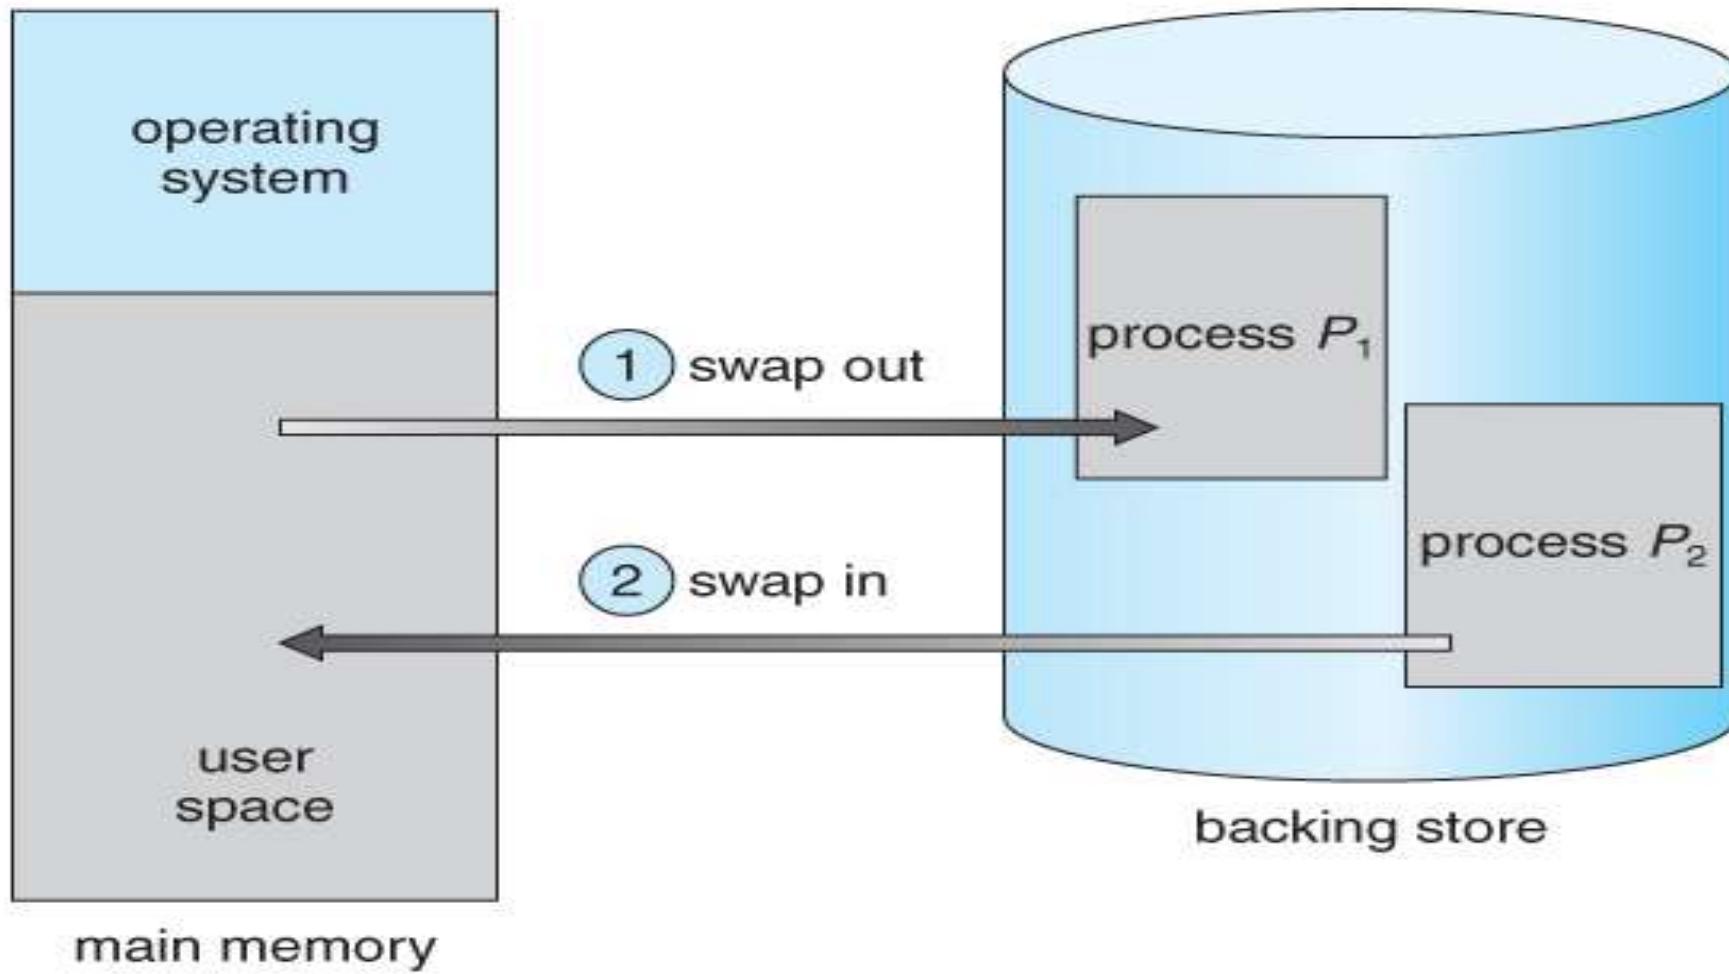
## Overlays:

- To enable a process to be **larger than** the amount of memory **allocated** to it
- Idea-
  - to keep in **memory only** those **instructions and data** **that are needed** at any given time.
  - When other instructions are needed, they are **loaded into** space occupied previously by **instructions** that are **no longer needed**.

## Swapping:

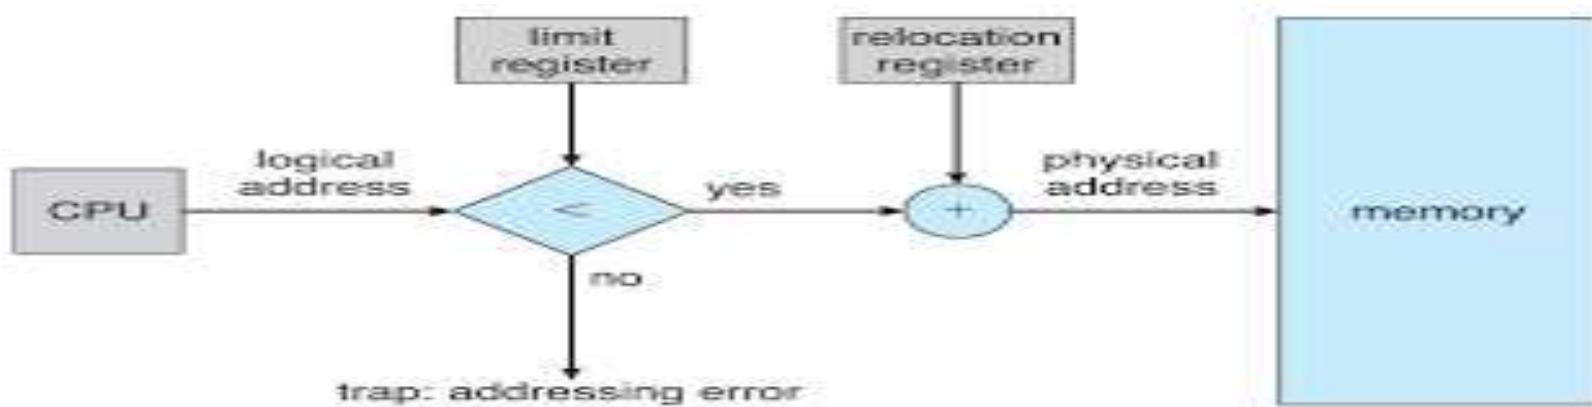
- A process needs to be in memory to be executed.
- A process, however, can be swapped temporarily out of memory to a backing store, and then brought back into memory for continued execution.
- A variant of this swapping policy is used for **priority-based scheduling algorithms**
  - When the higher-priority process finishes, the lower-priority process can be swapped back in and continued.
  - This variant of swapping is sometimes called **roll out, roll in**.

- The system maintains a ready queue consisting of all processes whose memory images are on the backing store or in memory and are ready to run.
- Whenever the CPU scheduler decides to execute a process, it calls the dispatcher.
- The dispatcher checks to see whether the next process in the queue is in memory.
- If not, and there is no free memory region, the dispatcher swaps out a process currently in memory and swaps in the desired process.
- It then reloads registers as normal and transfers control to the selected process.
- The context-switch time in such a swapping system is fairly high.
- To get an idea of the context-switch time,
  - let us assume that the user process is of size 1MB
  - the backing store is a standard hard disk with a transfer rate of 5 MB per second.
  - The actual transfer of the 1 MB process to or from memory takes  
 $1000 \text{ KB} / 5000 \text{ KB per second} = 1 / 5 \text{ second} = 200 \text{ milliseconds.}$



## Contiguous Memory Allocation:

- The memory is usually divided into **two partitions**:
  - one for the resident **operating system**, and one for the **user processes**.
  - We may place the operating system in either low memory or high memory.
- **contiguous** memory allocation, each **process** is contained in a **single contiguous section** of memory.
- Issue-protecting the operating system from user processes, and protecting user processes from one another.
- provide this protection by using a **relocation register**, with a **limit register**,



## Memory Allocation:

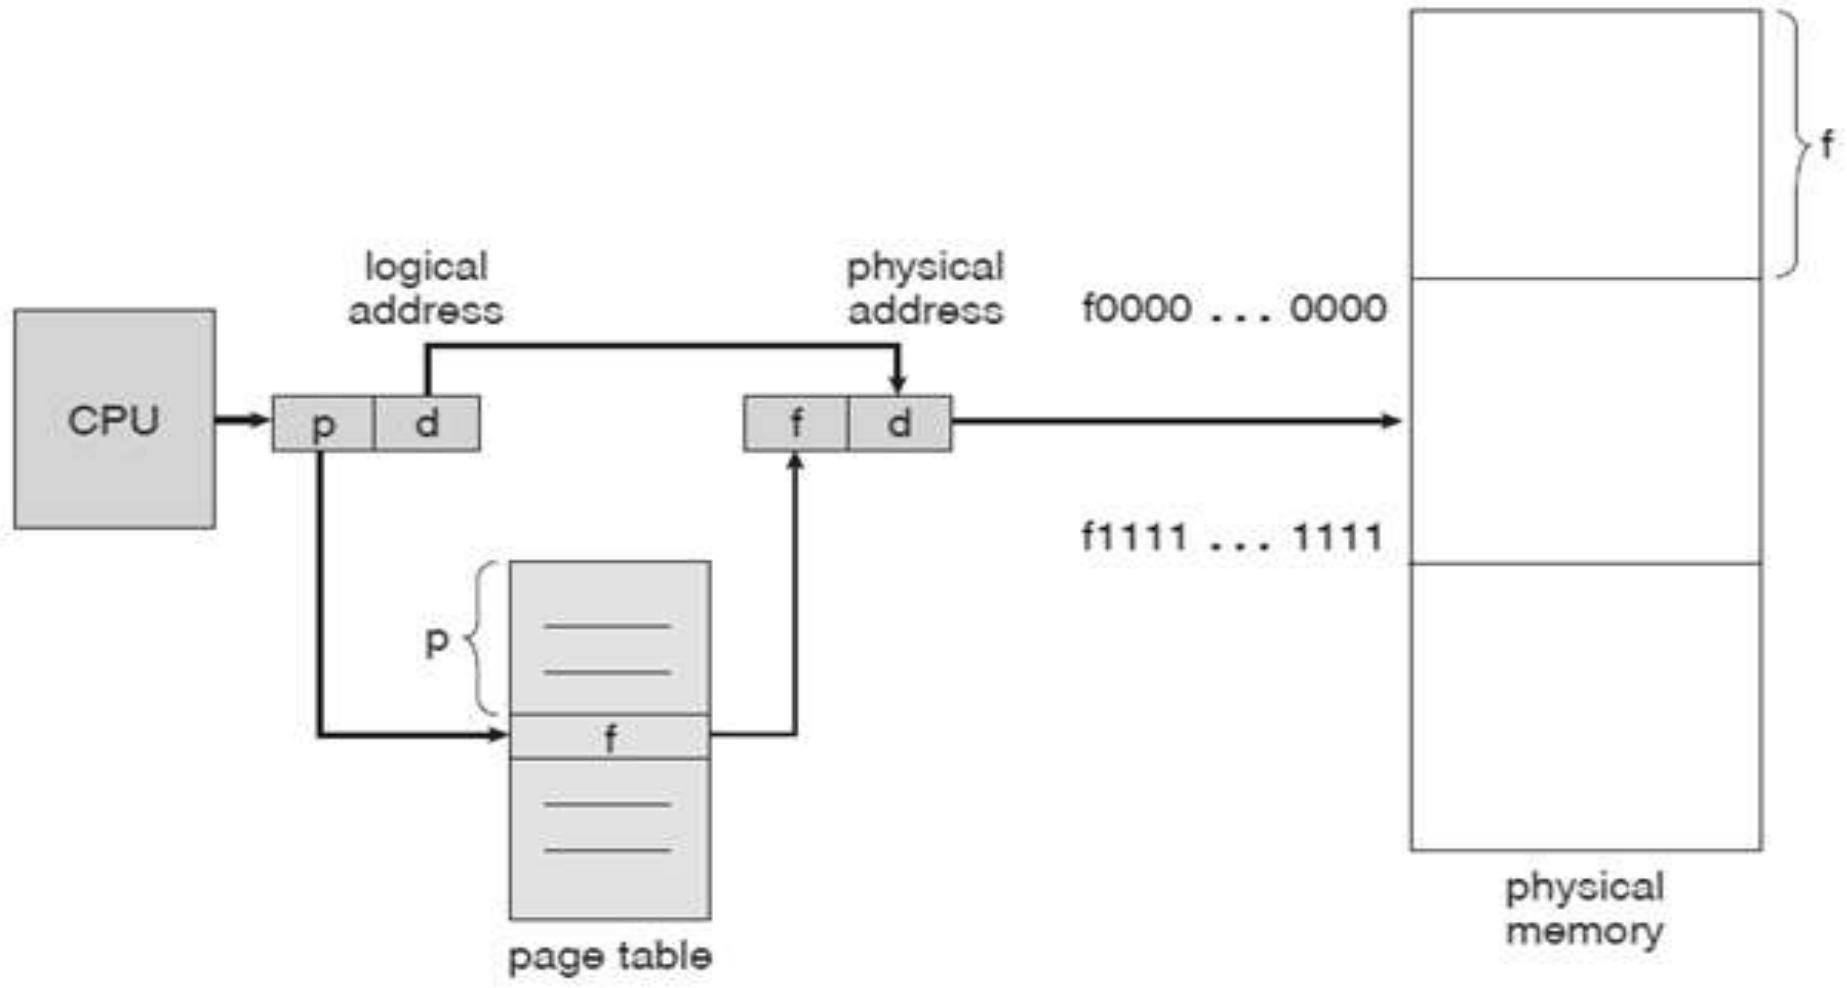
- simplest method for memory allocation (called **MFT**) is
  - to divide memory into **several fixed-sized partitions**.
  - Each partition may contain exactly one process
  - the degree of multiprogramming is bound by the number of partitions.
  - **it is no longer in use.**
- a **generalization** of the fixed-partition scheme (called **MVT**);
  - it is used primarily in a batch environment.
  - The operating system keeps a table indicating which parts of memory are available and which are occupied.
  - initially, all memory is available for user processes, and is considered as one **large block** of available memory, a **hole**.
  - When a **process** arrives and **needs memory**, we search for a hole large enough for this process.
  - **If we find** one, we **allocate** only as much memory as is needed,
  - **keeping the rest** available to satisfy **future** requests.
- how to satisfy a request of size **n** from a list of **free holes**.
  - **Firstfit**: Allocate the **first hole** that is **big enough**.
  - **Bestfit**: Allocate the **smallest hole** that is **big enough**.
  - **Worst fit**: Allocate the **largest hole** that is **big enough**.

## Fragmentation:

- The memory **allocated** to a process may be slightly **larger than the requested** memory.
- The **difference** between these two numbers is **internal fragmentation**-memory that is internal to a partition but is **not being used**.
- **External fragmentation** exists when enough **total memory space exists to satisfy a request, but it is not contiguous**;
  - storage is fragmented into a large number of small holes.
- One **solution** to the problem of external fragmentation is **compaction**.
- Compaction algorithm is simply to **move all processes** toward **one end of memory**; all holes move in the other direction, producing one large hole of available memory.
  - This scheme can be **expensive**.
- Another possible **solution** to the external-fragmentation problem is to permit the **logical address space** of a process to be **noncontiguous**,
- solution: **paging and segmentation**. These techniques can also be **combined**

## Paging:

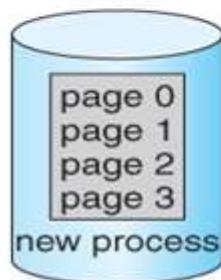
- Physical memory is broken into fixed-sized blocks called frames.
- Logical memory is also broken into blocks of the same size called pages.
- When a process is to be executed, its pages are loaded into any available memory frames from the backing store.
- The backing store is divided into fixed-sized blocks that are of the same size as the memory frames.
- Every address generated by the CPU is divided into two parts:  
a page number (p) and a page offset (d).
- The page number is used as an index into a page table.
- The page size (like the frame size) is defined by the hardware.
  - The size of a page is typically a power of 2, varying between 512 bytes and 16 MB per page, depending on the computer architecture.
- we have no external fragmentation: Any free frame can be allocated to a process that needs it. However, we may have some internal fragmentation.



- An important aspect of paging is the **clear separation between the user's view of memory and the actual physical memory**

free-frame list

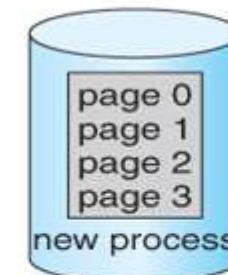
14  
13  
18  
20  
15



(a)

free-frame list

15



new-process page table

0	14
1	13
2	18
3	20

(b)

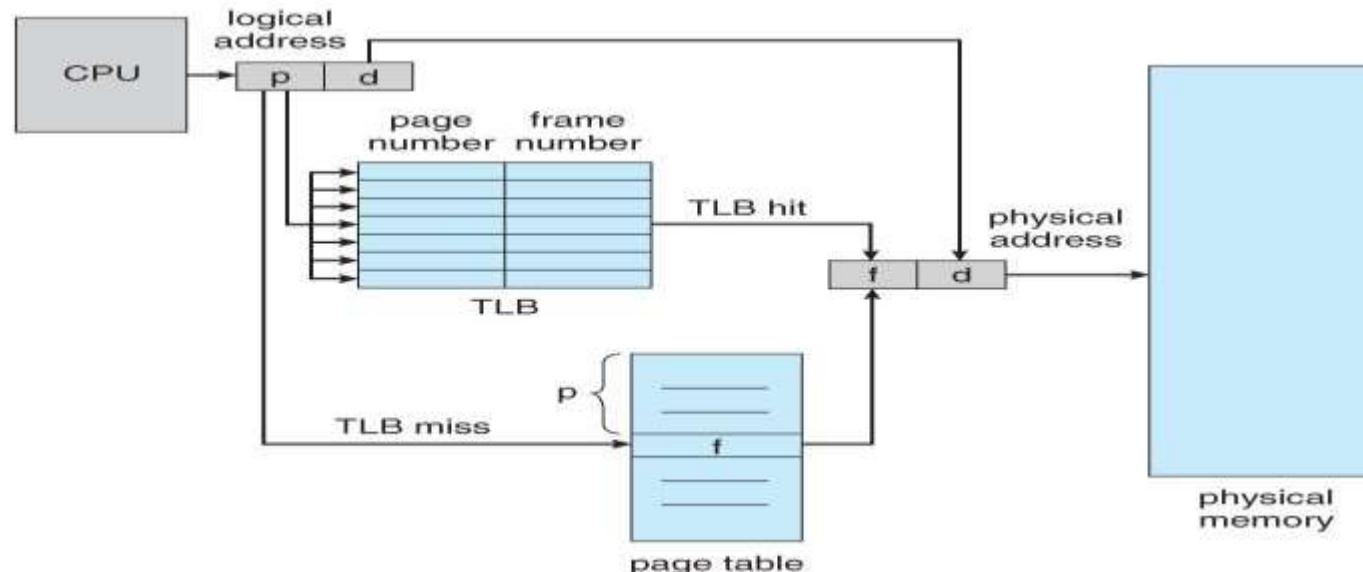
13 page 1  
14 page 0

15  
16  
17  
18 page 2  
19  
20 page 3  
21

## Hardware Support

to provide hardware support for this operation, in order to make it as fast as possible and to make process switches as fast as possible also

- to use a set of registers for the page table.
- to store the page table in main memory, and to use a single register ( called the **page-table base register, PTBR** ) to record where in memory the page table is locate
- to use a very special high-speed memory device called the **translation look-aside buffer, TLB**



The percentage of times that a particular page number is found in the TLB is called the **hit ratio**.

- An 80-percent hit ratio means that we find the desired page number in the TLB 80 percent of the time.
- If it takes 20 nanoseconds to search the TLB, and 100 nanoseconds to access memory,
- then a mapped memory access takes 120 nanoseconds when the page number is in the TLB

**Effective access Time=Hit time+miss time**

$$=(0.80 \times 120) + (0.20 \times 220)$$

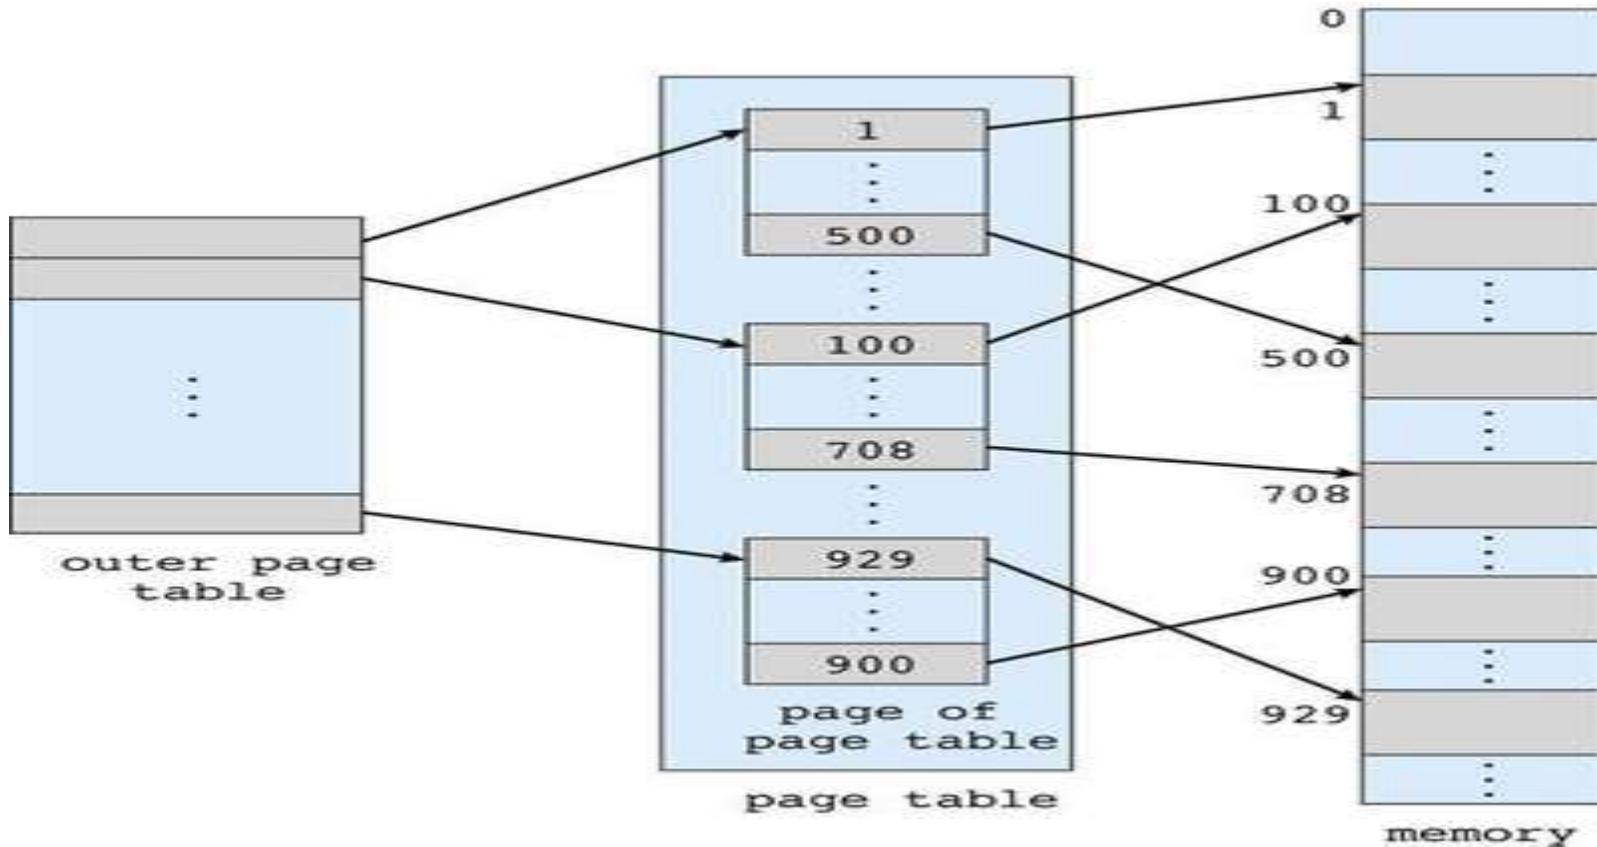
$$=140 \text{ nano seconds}$$

### **Protection:**

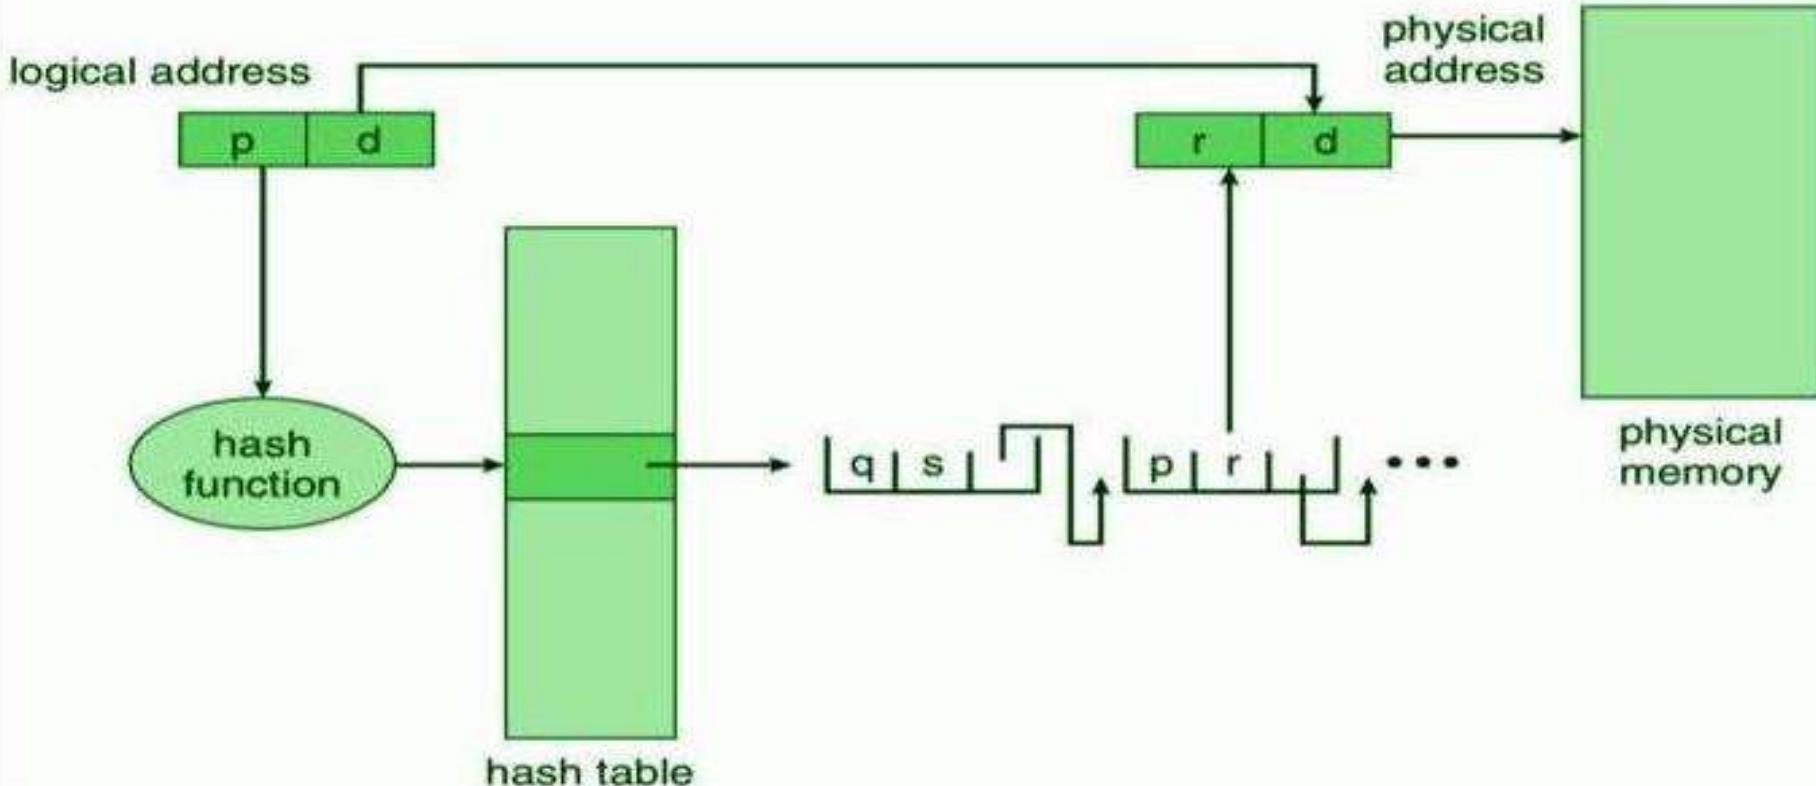
Memory protection in a paged environment is accomplished by protection bits that are associated with each frame. Normally, these bits are kept in the page table.

## Structure of the Page Table:

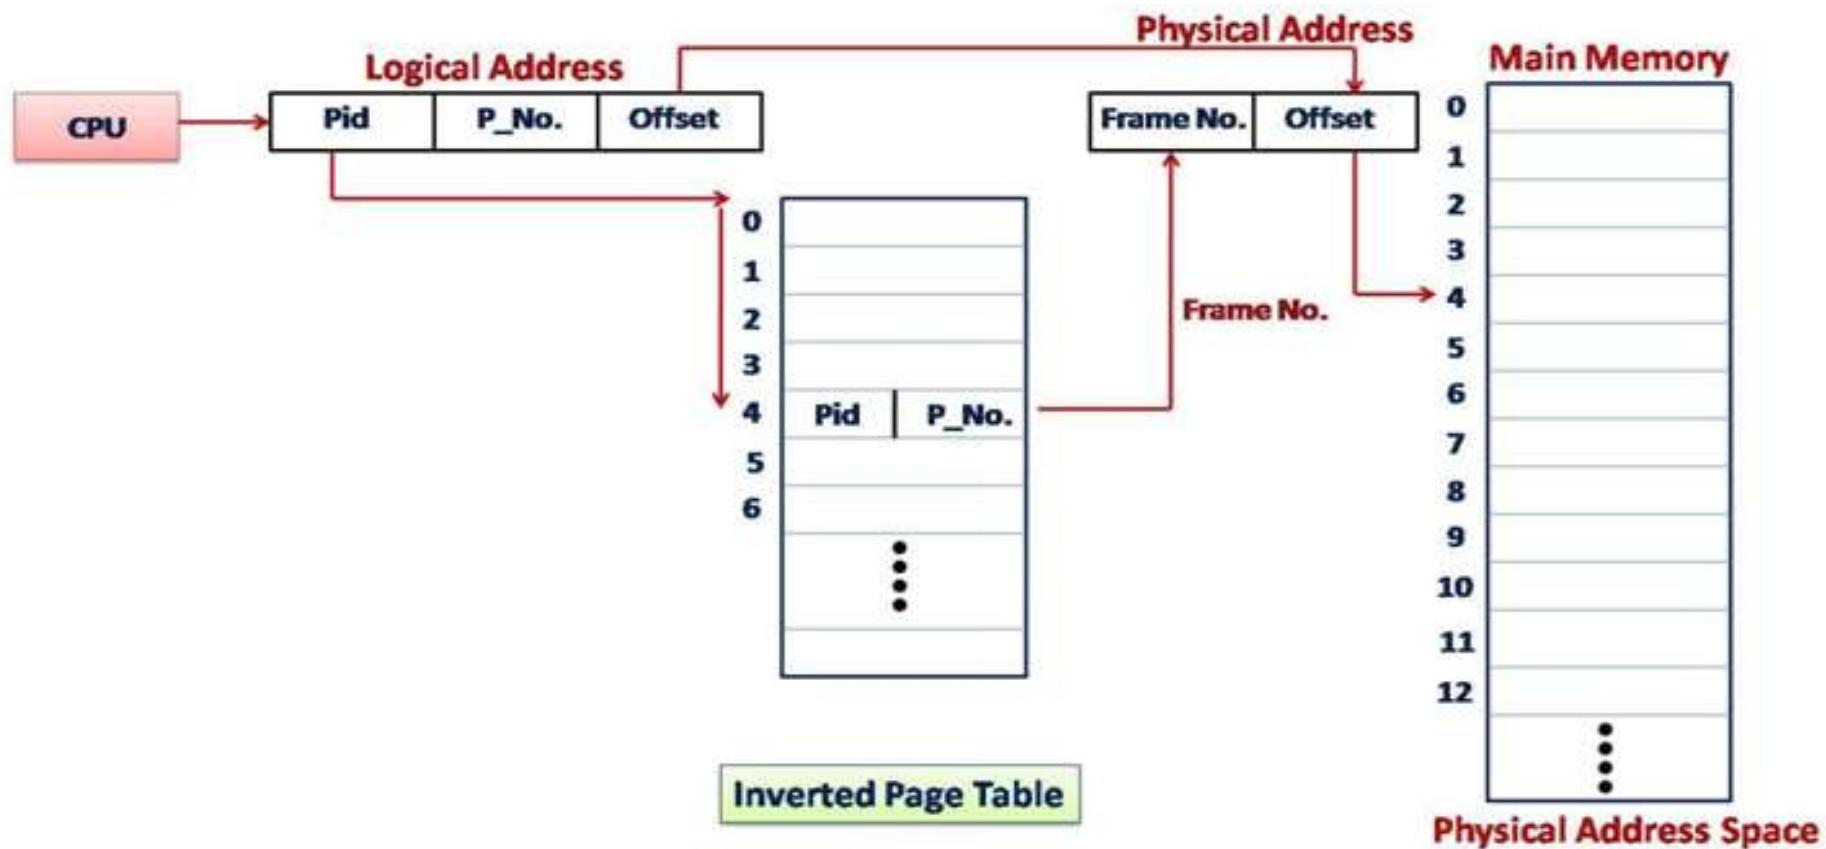
### Hierarchical Paging:



# Hashed Page Tables:

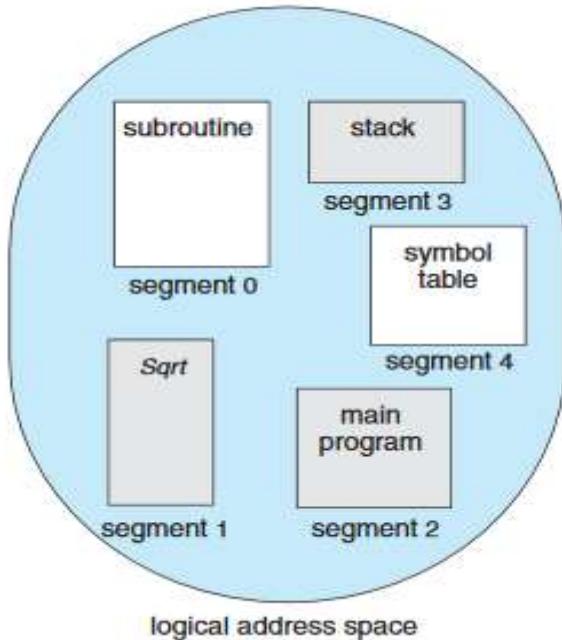


# Inverted Page Table:



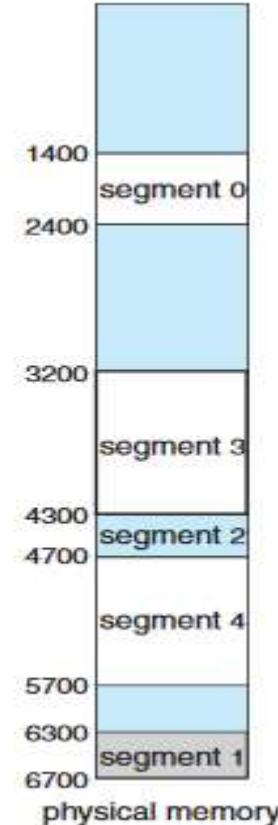
# Segmentation

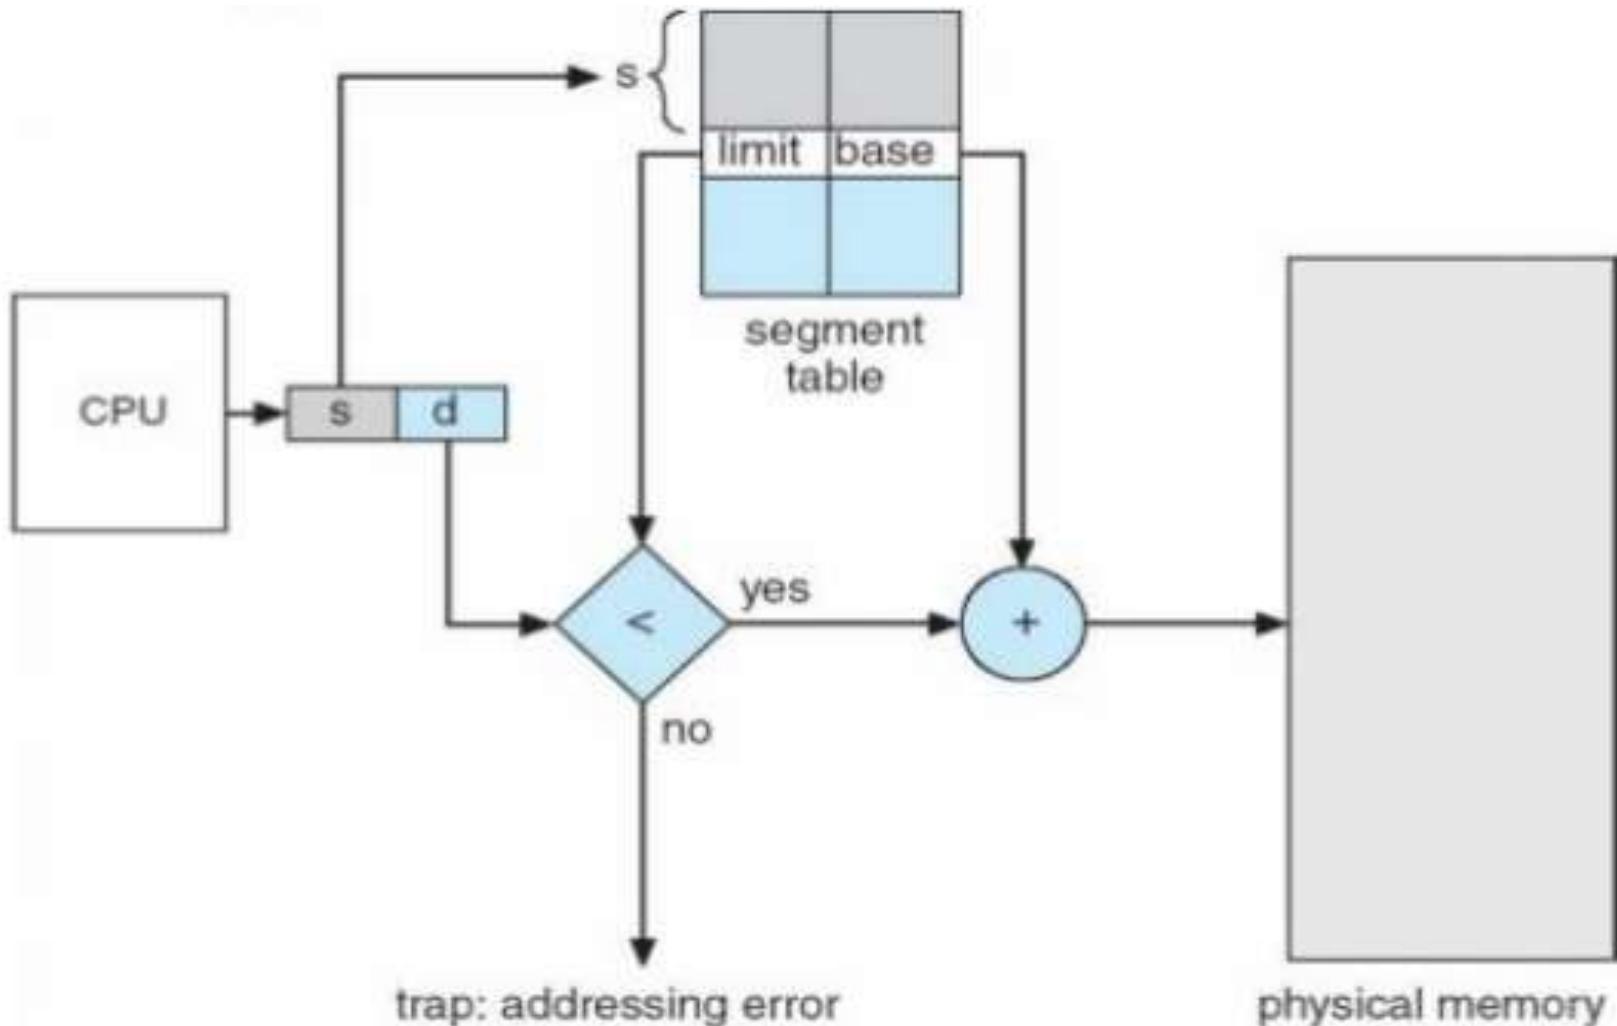
- paging is the separation of the user's view of memory and the actual physical memory.
- The user's view of memory is not the same as the actual physical memory.



	limit	base
0	1000	1400
1	400	6300
2	400	4300
3	1100	3200
4	1000	4700

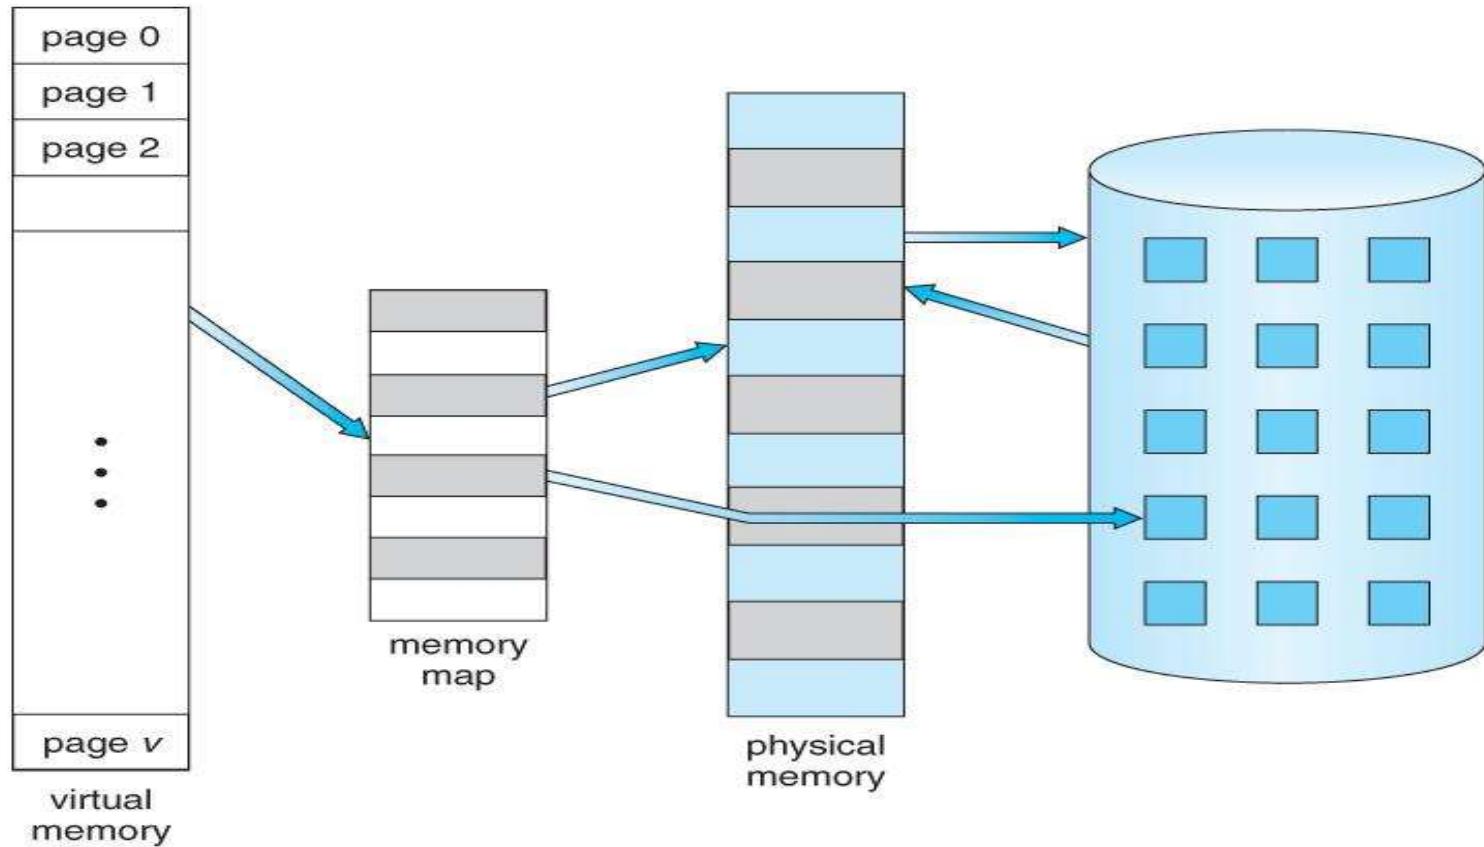
segment table





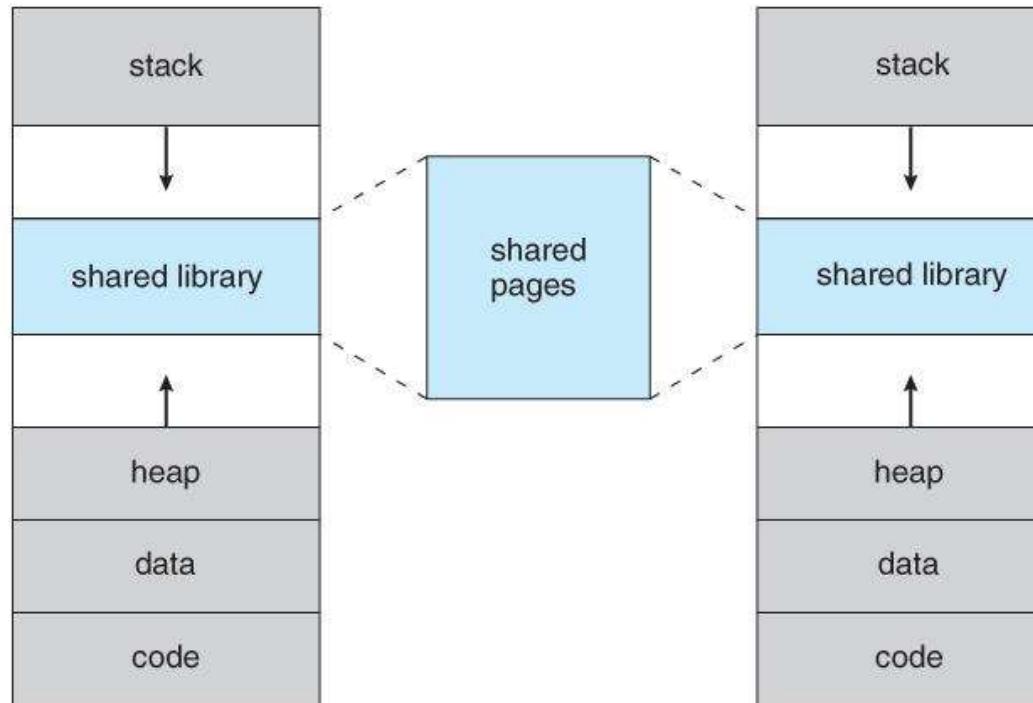
# Virtual memory

- is a technique that allows the execution of processes that may not be completely in memory



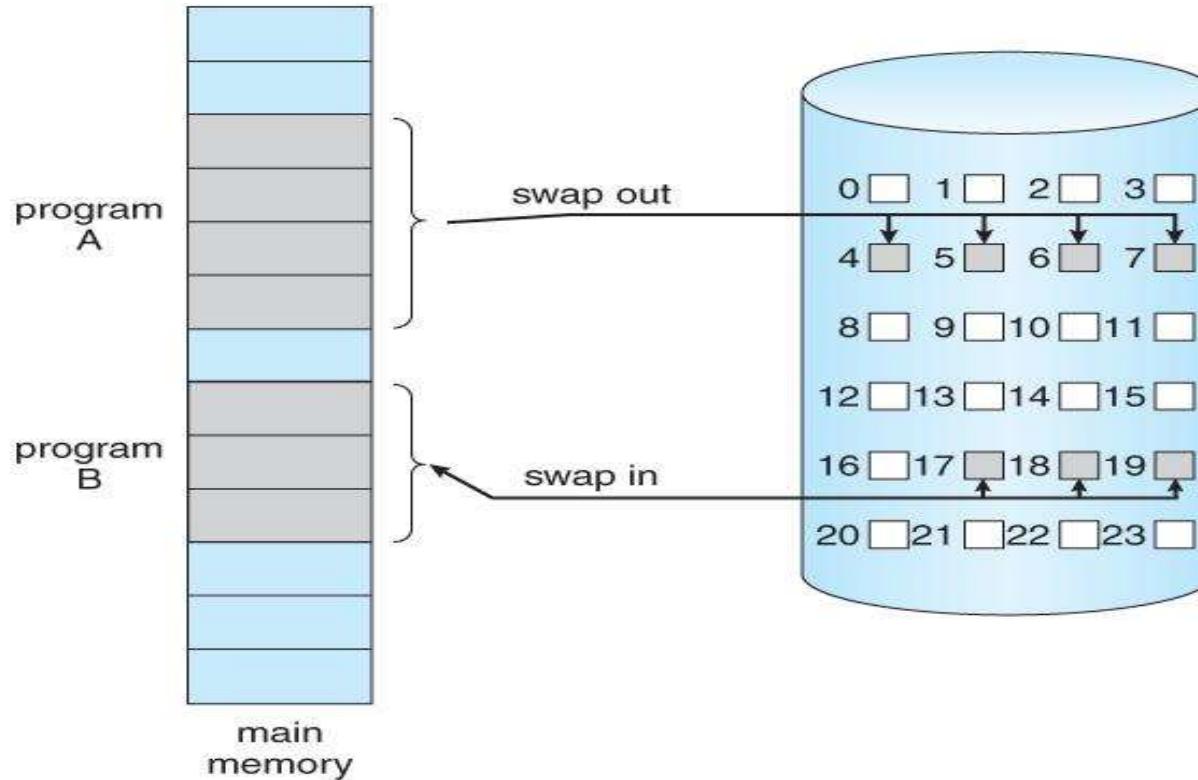
Virtual memory also allows the sharing of files and memory by multiple processes, with several benefits:

- System libraries
- Processes can also share virtual memory by mapping the same block of memory to more than one process.
- Process pages can be shared during a fork( ) system call,



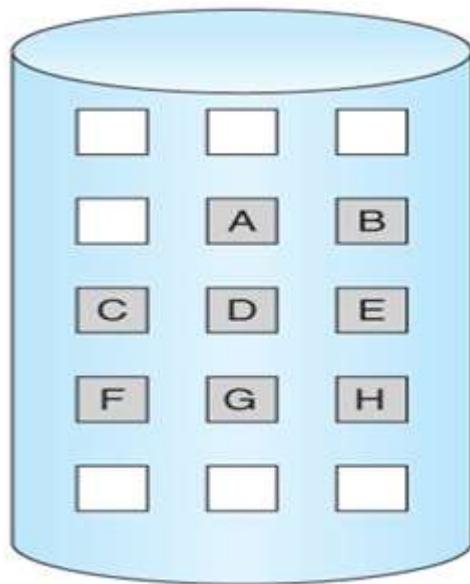
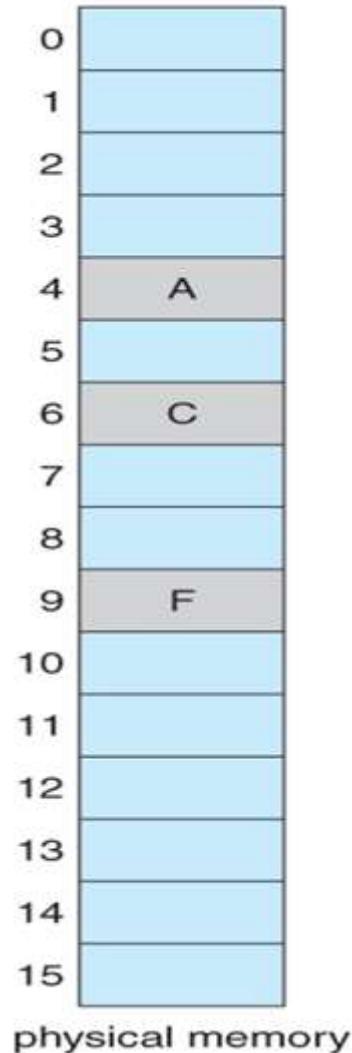
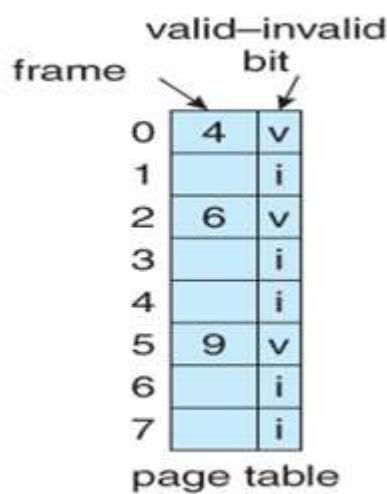
# Demand Paging

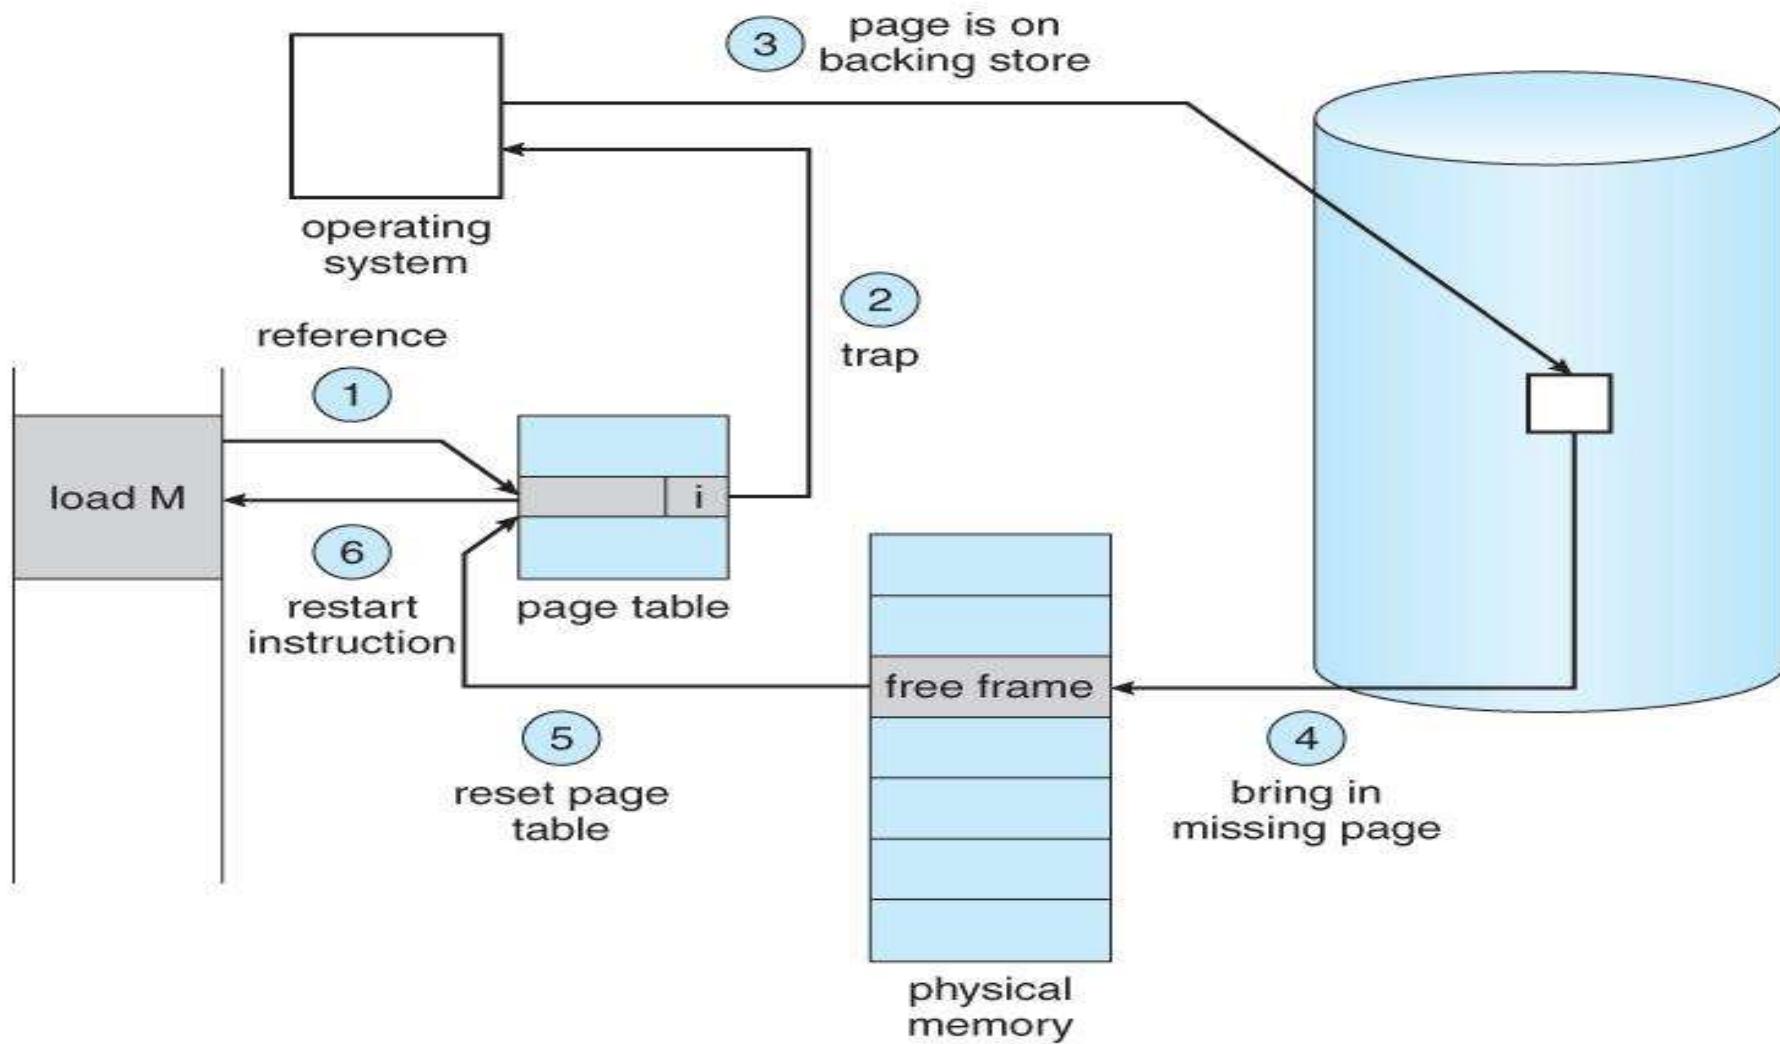
- when a process is swapped in, its pages are not swapped in all at once.
- Rather they are swapped in only when the process needs them. ( on demand. )
- This is termed a *lazy swapper*, although a *pager* is a more accurate term.



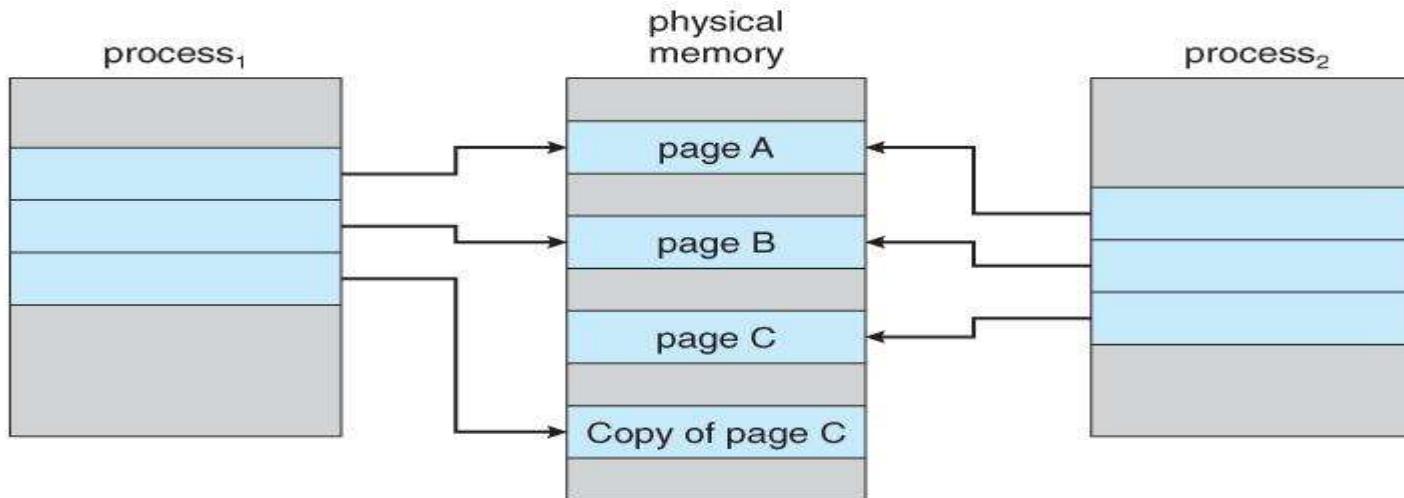
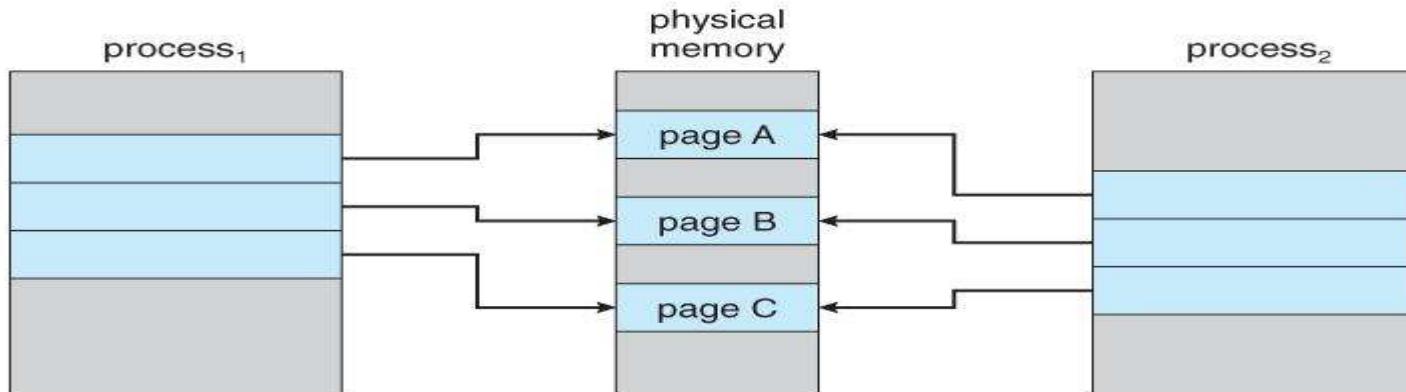
0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H

logical  
memory

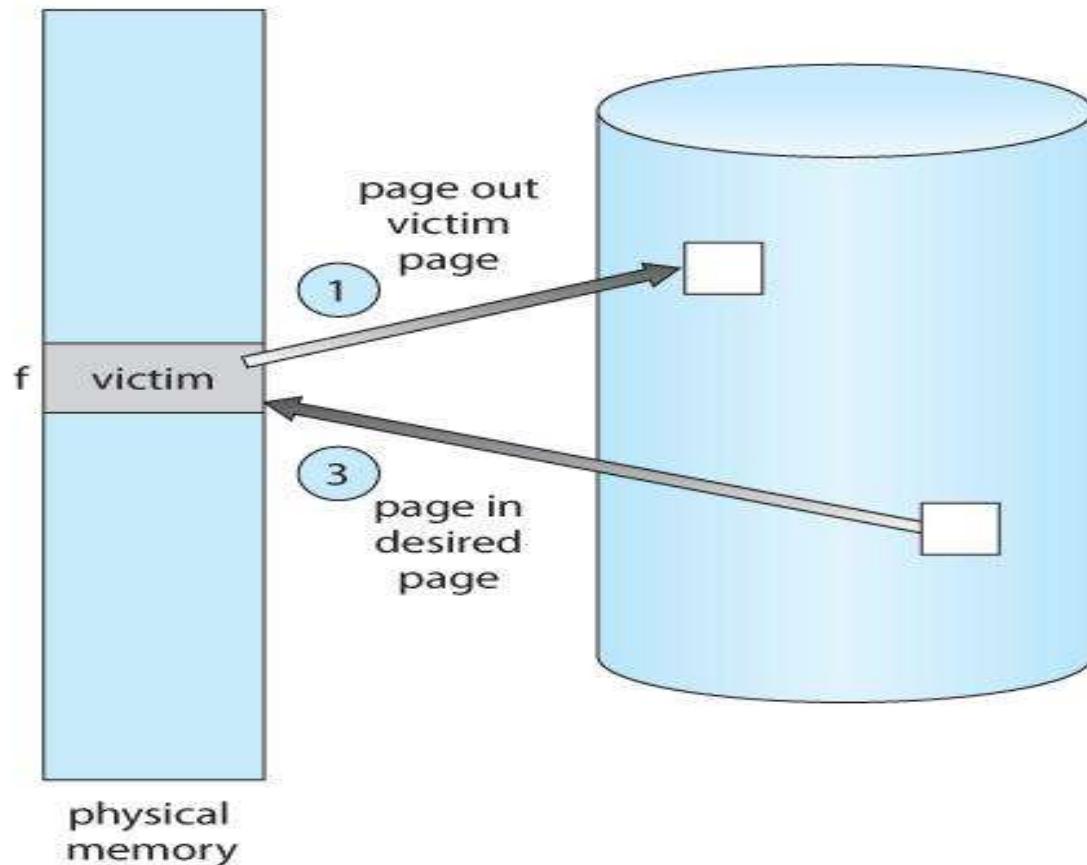
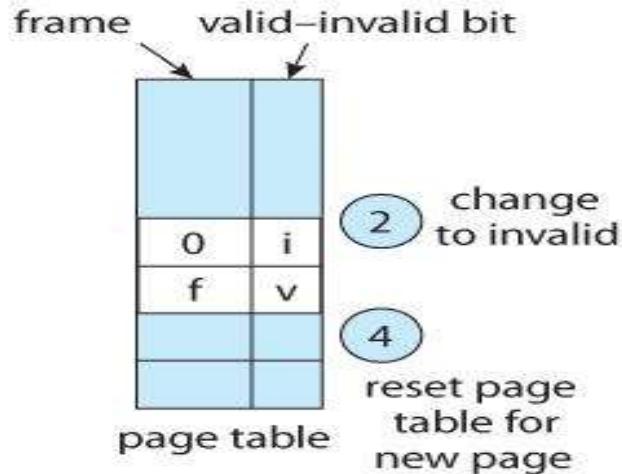




# Copy-on-Write



# Page Replacement



# Page replacement Algorithms

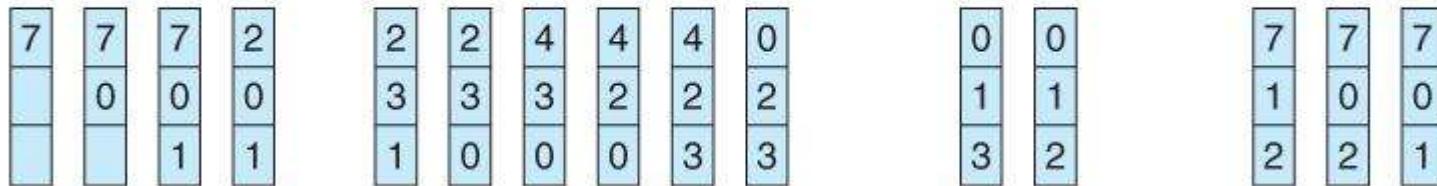
- FIFO Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

- *Belady's anomaly*

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

- **Optimal Page Replacement**

Replace the page that will not be used for the longest time in the future.

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

Reference string:

1 2 3 4 5 3 4 1 6 7 8 7 8 9 7 8 9 5 4 5 4 2

- LRU Page Replacement

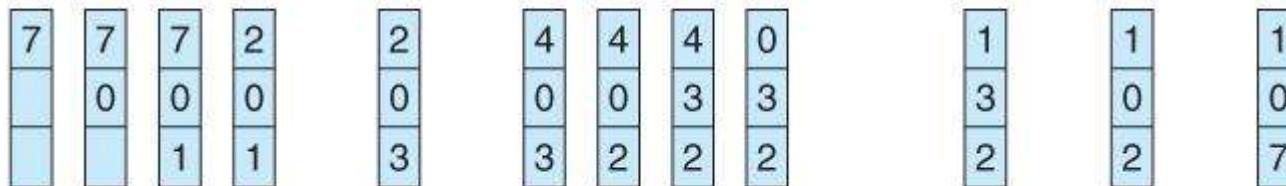
The page that has not been used in the longest time

reference string

7 0 1 2 0 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

reference string

4 7 0 7 1 0 1 2 1 2 7 1 2

- Two simple approaches commonly used

  - Counters
  - Stack

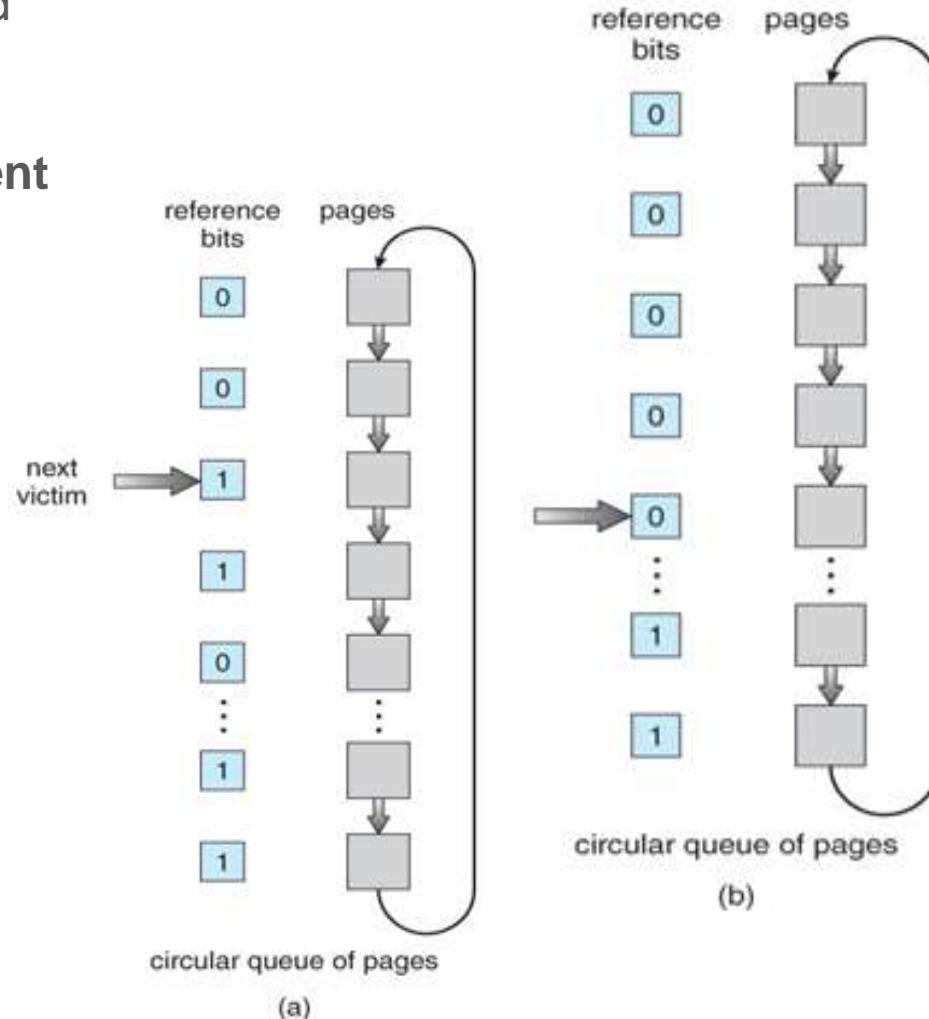
- **LRU-Approximation Page Replacement**

provide a *reference bit* for every entry in a page table

- **Additional-Reference-Bits Algorithm**

most recent 8 reference bits for each page

- **Second-Chance Algorithm**



- **Enhanced Second-Chance Algorithm**

- the reference bit and the modify bit ( dirty bit )
  - ( 0, 0 ) - Neither recently used nor modified.
  - ( 0, 1 ) - Not recently used, but modified.
  - ( 1, 0 ) - Recently used, but clean.
  - ( 1, 1 ) - Recently used and modified

- **Counting-Based Page Replacement :**

- *Least Frequently Used, LFU:*
- *Most Frequently Used, MFU:*

- **Page-Buffering Algorithms**

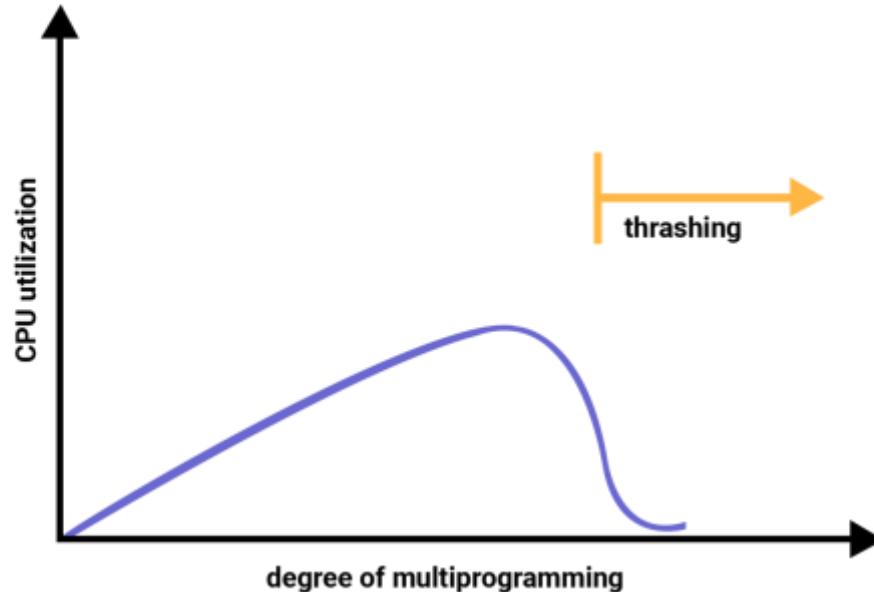
- **Frames Allocation:**

- minimum number of free frames at all times
- Equal Allocation
- Proportional Allocation

$$a_i = m * S_i / S$$

# Thrashing

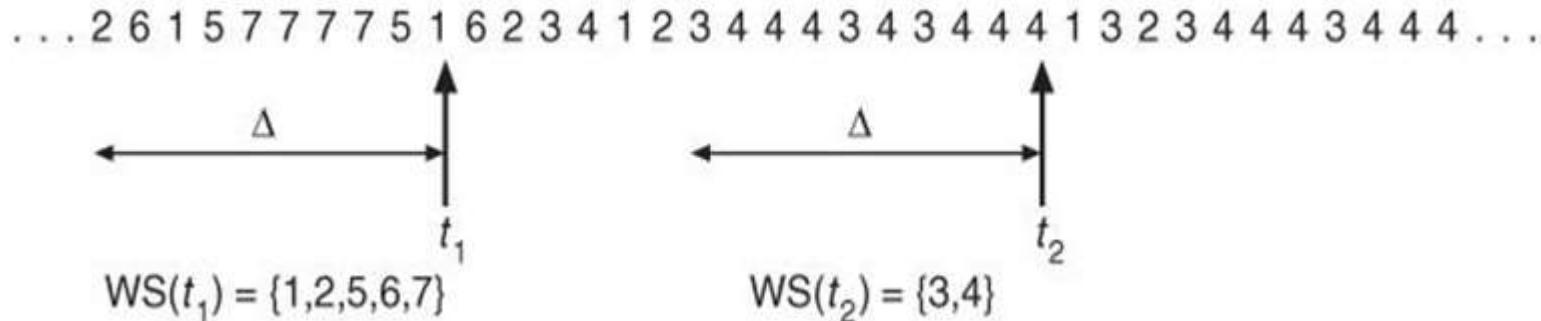
- A process is thrashing if it is spending more time paging than executing.
- limit the effects of thrashing by using a **local replacement algorithm** (or **priority replacement algorithm**).
- the **locality model of process execution**.



- Working-Set Model:

- uses a parameter,  $\Delta$ , to define the **working-set window**.
- **The idea is to examine the most recent  $\Delta$  page references.**
- **The set of pages in the most recent  $\Delta$  page references is the working set**

page reference table



- **Page-Fault Frequency:**

- upper and lower bounds on the desired page-fault rate
- If the actual page-fault rate exceeds the upper limit, we allocate that process another frame;
- if the page-fault rate falls below the lower limit, we remove a frame from that process.

