

Operating systems

By
I Ravindra kumar, B.Tech, M.Tech,(Ph.D.)

Assistant professor,
Dept of CSE, VNR VJIET

File System

File Concept:

- Computers can store information on several different storage media
 - magnetic disks, magnetic tapes, and optical disks
- Computer system will be convenient to use,
The operating system provides a uniform logical view of information storage.
- The operating system abstracts from the physical properties of its storage devices to define a logical storage unit (the file).
- Files are mapped, by the operating system, onto physical devices
- A file is a named collection of related information that is recorded on secondary storage

- **File Attributes**

A file has certain other attributes, which **vary from one operating system** to another,

Typically consist of these:

- **Name:**
- **Identifier:**
- **Type:**
- **Location:**
- **Size:**
- **Protection:**
- **Time, date, and user identification:**

- **File Operations:** A file is an **abstract data type**.
 - Creating a file:
 - Writing a file:
 - Reading a file:**current-file-position pointer**
 - Repositioning within a file:
 - Deleting a file:
 - Truncating a file:
- The operating system keeps a **small table** containing information **about all open files (the open-file table)**.
- several pieces of **information** are associated with an open file.
 - File pointer:
 - File open count:
 - Disk location of the file:
 - Access rights

File Types

file type	usual extension	function
executable	exe, com, bin or none	read to run machine-language program
object	obj, o	compiled, machine language, not linked
source code	c, cc, java, pas, asm, a	source code in various languages
batch	bat, sh	commands to the command interpreter
text	txt, doc	textual data, documents
word processor	wp, tex, rrf, doc	various word-processor formats
library	lib, a, so, dll, mpeg, mov, rm	libraries of routines for programmers
print or view	arc, zip, tar	ASCII or binary file in a format for printing or viewing
archive	arc, zip, tar	related files grouped into one file, sometimes compressed, for archiving or storage
multimedia	mpeg, mov, rm	binary file containing audio or A/V information

File Structure

- The operating system may
 - require that an executable file have a **specific structure**
 - it can determine **where in memory** to load the file
 - What is the location of the first instruction
- one of the disadvantages
 - multiple **file structures**:
- The resulting size of the operating system
 - **cumbersome**.
- If the operating system defines five different file structures,
 - it needs to contain the code to support these file structures.
- Severe problems may result from new applications that require information structured in ways not supported by the operating system.
- The Macintosh operating system also supports a minimal number of file structures.
- It expects files to contain two parts: **a resource fork and a data fork**.
- The resource fork contains information of interest to the user.

Internal File Structure

- Locating an offset within a file can be **complicated**
- All **disk I/O** is performed in units of one block (**physical record**), and all blocks are the **same size**.
 - Logical records may even vary in length.
- **Packing a number of logical records** into physical blocks is a common **solution** to this problem.
- For example, the **UNIX** operating system **defines all files** to be simply a **stream of bytes**.
 - Each byte is individually **addressable** by its offset from the **beginning** (or **end**) of the **file**. In this case, the logical record is **1 byte**.
- The **file system** automatically **packs and unpacks bytes** into **physical disk**
- The logical record size, physical block size, and packing technique determine how **many logical records are in each physical block**.
- The **packing** can be done either
 - the user's **application program** or
 - by the **operating system**.
- The wasted bytes allocated to keep everything in units of blocks (instead of bytes) is ***internal fragmentation***.
 - the **larger the block size**, the greater the **internal fragmentation**.

File-System Structure

- Two characteristics that make them a convenient medium for storing multiple files:
 - They can be rewritten in place;
 - They can access directly any given block of information on the disk.
- A file system poses two quite different design problems.
 - how the file system should look to the user
 - Creating algorithms and data structures to map the logical file system onto the physical secondary-storage devices.
- A file control block (**FCB**) contains
 - information about the file,
 - ownership,
 - permissions
 - location of the file contents
- UNIX uses the **UNIX file system (UFS) as a base.**
- Windows NT supports disk file-system formats of
 - **FAT, FAT32 and NTFS (or Windows NT File System),**
 - as well as CD-ROM, DVD, and floppy-disk file-system formats.




Figure 12.1 Layered file system.

- **Directory Implementation:**


selection of directory-allocation and directory-management algorithms has a large effect on the
 - efficiency, performance, and reliability of the file system
- **Linear List**
 - To use a linear list of file names with pointers to the data blocks.
 - Requires a linear search to find a particular entry.
 - Simple to program but time-consuming to execute

Hash Table

- a linear list stores the directory entries, but a hash data structure is also used.
- The hash table takes a value computed from the file name and returns a pointer to the file name in the linear list
- Insertion and deletion are also fairly straightforward,
- For example, assume that we make a linear-probing hash table that holds 64 entries.
- The hash function converts file names into integers from 0 to 63
- For instance, by using the remainder of a division by 64.
- If we later try to create a 65th file, we must enlarge the directory hash table-say, to 128 entries

Allocation Methods:

- how to **allocate space to these files** so that disk space is utilized effectively and files can be accessed quickly.
- **Contiguous Allocation:**
 - requires **each file to occupy a set of contiguous blocks on the disk**
 - Contiguous allocation of a file is **defined by the disk address and length (in block units) of the first block.**
 - If the file **is n blocks long and starts at location b ,** then it occupies blocks $b, b + 1, b + 2, \dots, b + n - 1$.
 - The **directory entry** for each file indicates the address of the **starting block and the length** of the area allocated for this file




Directory

file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2

- **Linked Allocation:**
 - With linked allocation, each file is a linked list of disk blocks;
 - the disk blocks may be scattered anywhere on the disk.
 - The directory contains a pointer to the first and last blocks of the file.
 - An important variation on the linked allocation method is the use of a file allocation table (**FAT**).
 - This simple but efficient method of disk-space allocation
 - It is used by the MS-DOS and OS/2 operating systems.
 - A section of disk at the beginning of each partition is set aside to contain the table.

Directory


file	start	end
jeep	9	25



- **Indexed Allocation**
 - the **pointers to the blocks** are scattered with the blocks themselves all over the disk and need to be retrieved
 - in order, **Indexed allocation solves this problem by bringing all the pointers together** into one location: the **index block**.
 - Each file has its **own index block**, which is an array of disk-block addresses.
 - The **ith entry** in the index block points to the **ith block** of the file.
 - The directory contains the **address of the index block**

Directory

file	index block
jeep	19




If the **index block** is too small, however, it will not be able to hold enough pointers for a **large file**, and a mechanism will have to be available to deal with this issue:

- **Linked scheme:**
- **Multilevel index:**
- **Combined scheme:**
 - say, **15 pointers** of the **index block** in the file's inode.
 - The **first 12** of these pointers point to **direct blocks**;
 - The next **3 pointers** point to **indirect blocks**.
 - **The first indirect block pointer** is the address of a **single indirect block**.
 - Then there is a **double indirect block pointer**,
 - The last pointer would contain the address of a **triple indirect block**.

- **Free-Space Management:**
 - keep track of free disk space, the system maintains a **free-space list**.
 - **The free-space list records all** free disk blocks-those not allocated to some file or directory.
- **Bit Vector:**
 - free-space list is implemented as a bit **map or bit vector**
 - Each block is represented by 1 bit. If the **block is free**, the **bit is 1**; if the block is **allocated**, the bit is **0**.
 - its relatively **simplicity and efficiency** in finding the **first free block, or *n consecutive free blocks on the disk***.
 - The first **non-0 word** is scanned for the **first 1 bit**, which is the **location of the first free block**.
 - The calculation of the block number is
(number of bits per word) x (number of 0-value words) + offset of first 1 bit.

- **Linked List:**
 - keeping a pointer to the **first free block** in a **special location** on the disk and caching it in memory




- **Grouping:**
 - to store the addresses of n free blocks in the first free block.
 - The first $n-1$ of these blocks are actually free.
 - The last block contains the addresses of another n free blocks
- **Counting:**
 - Keep the address of the first free block and the number n of free contiguous blocks that follow the first block.
 - Each entry in the free-space list then consists of a disk address and a count.

Efficiency:


- For instance, **UNIX inodes** are preallocated on a partition.
- Even an "empty" disk has a percentage of its space lost to inodes.
- by preallocating the inodes and spreading them across the partition, we improve the file system's performance
- BSD UNIX varies the cluster size as a file grows.
- Large clusters are used where they can be filled, and small clusters are used for small files and the last cluster of a file.
- The types of data normally kept in a file's directory (or inode) entry also require consideration.
- Commonly, a "last write date" is recorded to supply information to the user and to determine whether the file needs to be backed up.
- Some systems also keep a "last access date," so that a user can determine when the file was last read.
- The result of keeping this information is that, whenever the file is read, a field in the directory structure must be written to.
- This change requires the block to be read into memory, a section changed, and the block written back out to disk, because operations on disks occur only in block (or cluster) chunks.

Performance:

- disk controllers include local memory to form an on-board **cache** that is sufficiently large to store entire tracks at a time
- maintain a separate section of main memory for a **disk cache**, where blocks are kept under the assumption that they will be used again shortly
- Other systems cache file data using a **page cache**



- Some versions of UNIX provide a **unified buffer cache**
- Consider the two alternatives of opening and accessing a file.
 - One approach is **to use memory mapping**
 - the second is to use the **standard system calls read and write**



- **Synchronous writes**
 - occur in the order in which the disk subsystem receives them, and the writes are not buffered.
 - Thus the calling routine must wait for the data to reach the disk drive before it can proceed.
- **Asynchronous writes**
 - are done the majority of the time
- **Free-behind**
 - removes a page from the buffer as soon as the next page is requested.
 - The previous pages are not likely to be used again and waste buffer space.
- **read-ahead,**
 - a requested page and several subsequent pages are read and cached